



# Benchmarking Concurrent Programs

Scott West

ETH Zürich

November 23, 2012



# MEASURE, THEN CUT

One result from the parallel language study: SCOOP has poor performance on shared-memory benchmarks.

A successful application of benchmarking: we found a problem!



# MEASURE, THEN CUT

One result from the parallel language study: SCOOP has poor performance on shared-memory benchmarks.

A successful application of benchmarking: we found a problem!



# CONCURRENCY V. PARALLELISM

## Definition (Parallelism)

A program doing more than one thing at a time to compute a result faster. The result is deterministic.

## Definition (Concurrency)

A program with multiple threads of control that which interleave in a non-deterministic way.

They are not exclusive: concurrency can be used to implement parallelism.



## OUR TASK

A performance benchmark should be a way to measure a relevant aspect of execution properties.

Parallelism benchmarks define a task that should have the same result every time, and the result must arrive quickly!

We aim to define a way to benchmark concurrent programs.



# OUR TASK

A performance benchmark should be a way to measure a relevant aspect of execution properties.

Parallelism benchmarks define a task that should have the same result every time, and the result must arrive quickly!

We aim to define a way to benchmark concurrent programs.



## OUR TASK

A performance benchmark should be a way to measure a relevant aspect of execution properties.

Parallelism benchmarks define a task that should have the same result every time, and the result must arrive quickly!

We aim to define a way to benchmark concurrent programs.



# WHAT IS A CONCURRENT BENCHMARK?

Since a concurrent program executes nondeterministically, benchmarking will concern profiling the performance characteristics of various *synchronization*, *communication*, and *coordination* tasks.





# BENCHMARK OVERVIEW

We wish to examine *scalability*: how does a benchmark react to increased contention/threads/problem size.

We will examine concurrent programs with multiple patterns:

**no share** threads do not share memory or coordinate.

**1-N comm.** a single master thread dispatches work to multiple workers.

**M-N comm.** multiple master threads dispatch work

**mutex** multiple threads contend for a critical section with high contention.

**benign share** threads share memory, but do not have guarded access to it.



# NONINTERFERENCE

## Purpose

*To test the cost of running multiple threads by the runtime system.*

Some concurrency abstractions have non-trivial departures from POSIX-like threads (such as green threads) and this such a test would determine the effect of any such abstractions.

This would take the shape of running some compute heavy task in each thread. I.e.

```
fib (30)
```



# 1-N COMMUNICATION

## Purpose

*To examine the cost of having multiple threads contending for a single event.*

This test replicates a common design in server applications where a main thread dispatches work to multiple reader threads.

The test will have a single channel of communication where the master dispatches integers to the workers. The workers continually receive and discard the integers.



# M-N COMMUNICATION

## Purpose

*To test the effect of multiple producers accessing a single consumption channel.*

Similar to the previous test, but with multiple masters and workers at either end of the work distribution.



# MUTEX

## Purpose

*To quantify the cost of gaining exclusive access to a critical section.*

The critical section should be heavily contended as this is the case where performance issues tend to arise.

The mutex operations may be noops (if a CS is not explicitly required), but *must* update some shared state in a simple way (ie, adding 1).



# INTERFERING

## Purpose

*To examine the overhead of continually accessing a shared resource without any extra protection not provided by the implementation.*

Very often there are benign dataraces, this test quantifies the performance cost of modifying shared memory without an explicit mutex protecting it.



# PRELIMINARY RESULTS

Benchmark	STM (Haskell)	SCOOP
No share	—	—
1-N comm.	—	—
M-N comm.	0.478 s	63.747 s
Mutex	0.067 s	24.235 s
Benign share	—	—



## GOING FURTHER

- ▶ Add more microbenchmarks, gathered from real situations.
- ▶ Implement the benchmarks in various concurrency approaches, with a focus on languages that focus on concurrency rather than data parallelism.
  - ▶ Candidates for other approaches: Erlang, Go, Ada, Clojure, F#, D, Java, Scala, Orc...
- ▶ Validation step: use them to predict performance of larger programs.





# PUBLICATION

Rough plan: ISSTA submission early next year.