



SCOOP performance and implementation

Scott West

ETH Zürich

February 19, 2013



EXAMINING SCOOP PERFORMANCE

Since the research retreat:

- ▶ Added the missing benchmarks that were discussed
- ▶ Added 2 more languages: Scala (Akka) actors and Java (mutex/condition variables)
- ▶ Started to investigate SCOOP performance and solutions



EXAMINING SCOOP PERFORMANCE

Since the research retreat:

- ▶ Added the missing benchmarks that were discussed
- ▶ Added 2 more languages: Scala (Akka) actors and Java (mutex/condition variables)
- ▶ Started to investigate SCOOP performance and solutions



EXAMINING SCOOP PERFORMANCE

Since the research retreat:

- ▶ Added the missing benchmarks that were discussed
- ▶ Added 2 more languages: Scala (Akka) actors and Java (mutex/condition variables)
- ▶ Started to investigate SCOOP performance and solutions

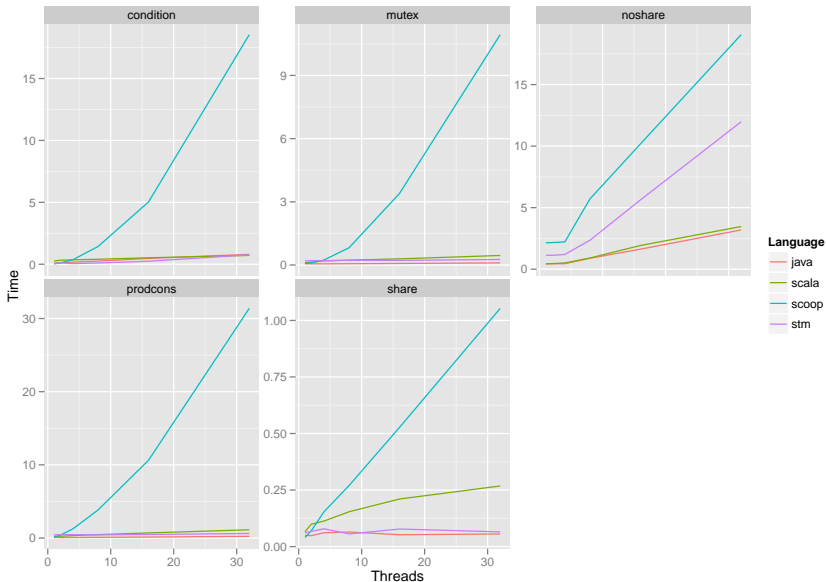


BENCHMARK OVERVIEW

- condition** two groups of threads, each increments a shared variable when it is odd/even.
- mutex** multiple threads contend for a critical section.
- no share** threads do not share memory or coordinate, perform a computationally bound task.
- prodcons** multiple producers and consumers work on a single queue.
- share** multiple threads all threads share memory, but do not guard access to it (if possible)



BENCHMARK RESULTS





NEXT STEPS

Theme from last time: measure then cut.

- ▶ SCOOP currently slows down superlinearly in the amount of work :(
 - ▶ How do we find what we should cut?
 - ▶ Profiling and code inspection.



NEXT STEPS

Theme from last time: measure then cut.

- ▶ SCOOP currently slows down superlinearly in the amount of work :(
:(
- ▶ How do we find what we should cut?
- ▶ Profiling and code inspection.



NEXT STEPS

Theme from last time: measure then cut.

- ▶ SCOOP currently slows down superlinearly in the amount of work :(
:(
- ▶ How do we find what we should cut?
- ▶ Profiling and code inspection.



SCOOP RUNTIME: AIMED AT SPEED

- ▶ Lots of atomic operations (compare and swap, atomic increment, etc.)
- ▶ Implemented records using arrays and offset indices to avoid the object overhead.
- ▶ Avoids synchronization primitives (locks, semaphores).



RESOURCE ACQUISITION

```
from
  acquired := False
until
  acquired
loop
  old_val :=
    cas (resrc, free, this)
  if old_val = free then
    acquired := True
  else
    yield
  end
end
```

- ▶ This code avoids a mutex!
- ▶ Mutexes are slow because they may trigger context switch...
- ▶ ... but *yield* also triggers a context switch.
- ▶ This code is basically a mutex but never sleeps and eats a core (it spins until it acquires the resource).
- ▶ Replacing it with a traditional mutex helps.



RESOURCE ACQUISITION

```
from
    acquired := False
until
    acquired
loop
    old_val :=
        cas (resrc, free, this)
    if old_val = free then
        acquired := True
    else
        yield
    end
end
```

- ▶ This code avoids a mutex!
- ▶ Mutexes are slow because they may trigger context switch...
- ▶ ... but *yield* also triggers a context switch.
- ▶ This code is basically a mutex but never sleeps and eats a core (it spins until it acquires the resource).
- ▶ Replacing it with a traditional mutex helps.



RESOURCE ACQUISITION

```
from
  acquired := False
until
  acquired
loop
  old_val :=
    cas (resrc, free, this)
  if old_val = free then
    acquired := True
  else
    yield
  end
end
```

- ▶ This code avoids a mutex!
- ▶ Mutexes are slow because they may trigger context switch...
- ▶ ... but *yield* also triggers a context switch.
- ▶ This code is basically a mutex but never sleeps and eats a core (it spins until it acquires the resource).
- ▶ Replacing it with a traditional mutex helps.



RESOURCE ACQUISITION

```
from
  acquired := False
until
  acquired
loop
  old_val :=
    cas (resrc, free, this)
  if old_val = free then
    acquired := True
  else
    yield
  end
end
```

- ▶ This code avoids a mutex!
- ▶ Mutexes are slow because they may trigger context switch...
- ▶ ... but *yield* also triggers a context switch.
- ▶ This code is basically a mutex but never sleeps and eats a core (it spins until it acquires the resource).
- ▶ Replacing it with a traditional mutex helps.



RESOURCE ACQUISITION

```
from
  acquired := False
until
  acquired
loop
  old_val :=
    cas (resrc, free, this)
  if old_val = free then
    acquired := True
  else
    yield
  end
end
```

- ▶ This code avoids a mutex!
- ▶ Mutexes are slow because they may trigger context switch...
- ▶ ... but *yield* also triggers a context switch.
- ▶ This code is basically a mutex but never sleeps and eats a core (it spins until it acquires the resource).
- ▶ Replacing it with a traditional mutex helps.



MOVING FORWARD

- ▶ I started to rewrite the slow parts as best I could; these did lead to some improvements.
- ▶ However the changes are untestable: the code currently deadlocks and has some other bugs so I have no way to figure out if it's more or less broken after my changes.
- ▶ Develop a “laboratory” where I can try out ideas: proof-of-concept in C++.



MOVING FORWARD

- ▶ I started to rewrite the slow parts as best I could; these did lead to some improvements.
- ▶ However the changes are untestable: the code currently deadlocks and has some other bugs so I have no way to figure out if it's more or less broken after my changes.
- ▶ Develop a “laboratory” where I can try out ideas: proof-of-concept in C++.



MOVING FORWARD

- ▶ I started to rewrite the slow parts as best I could; these did lead to some improvements.
- ▶ However the changes are untestable: the code currently deadlocks and has some other bugs so I have no way to figure out if it's more or less broken after my changes.
- ▶ Develop a “laboratory” where I can try out ideas: proof-of-concept in C++.



SCOOP IMPLEMENTATION MODEL

A SCOOP processor is basically a work-queue which other processors have exclusive access to.



Naïve implementation STL queue wrapped in a monitor: 404s

Non-blocking TBB queue: 180s

Blocking TBB queue: 108s

Is this any faster than what we have now?

Sort of, EiffelStudio 7.1: 33s but EiffelStudio 7.2: 155s.



SCOOP IMPLEMENTATION MODEL

A SCOOP processor is basically a work-queue which other processors have exclusive access to.



Naïve implementation STL queue wrapped in a monitor: 404s

Non-blocking TBB queue: 180s

Blocking TBB queue: 108s

Is this any faster than what we have now?

Sort of, EiffelStudio 7.1: 33s but EiffelStudio 7.2: 155s.



SCOOP IMPLEMENTATION MODEL

A SCOOP processor is basically a work-queue which other processors have exclusive access to.



Naïve implementation STL queue wrapped in a monitor: 404s

Non-blocking TBB queue: 180s

Blocking TBB queue: 108s


Is this any faster than what we have now?

Sort of, EiffelStudio 7.1: 33s but EiffelStudio 7.2: 155s.



SCOOP IMPLEMENTATION MODEL

A SCOOP processor is basically a work-queue which other processors have exclusive access to.



Naïve implementation STL queue wrapped in a monitor: 404s

Non-blocking TBB queue: 180s

Blocking TBB queue: 108s

Is this any faster than what we have now?

Sort of, EiffelStudio 7.1: 33s but EiffelStudio 7.2: 155s.



SCOOP IMPLEMENTATION MODEL

A SCOOP processor is basically a work-queue which other processors have exclusive access to.



Naïve implementation STL queue wrapped in a monitor: 404s

Non-blocking TBB queue: 180s

Blocking TBB queue: 108s

Is this any faster than what we have now?

Sort of, EiffelStudio 7.1: 33s but EiffelStudio 7.2: 155s.



SCOOP IMPLEMENTATION MODEL

A SCOOP processor is basically a work-queue which other processors have exclusive access to.



Naïve implementation STL queue wrapped in a monitor: 404s

Non-blocking TBB queue: 180s

Blocking TBB queue: 108s


Is this any faster than what we have now?

Sort of, EiffelStudio 7.1: 33s but EiffelStudio 7.2: 155s.



SCOOP IMPLEMENTATION MODEL 2

Lock contention can be reduced if there is no memory to contend over:
use a queue of queues



Queue of queues: 15s.


With limited dequeue spinning: 3.7s

Don't use the queue for "second" queries: 2.0s



SCOOP IMPLEMENTATION MODEL 2

Lock contention can be reduced if there is no memory to contend over:
use a queue of queues



Queue of queues: 15s.


With limited dequeue spinning: 3.7s

Don't use the queue for "second" queries: 2.0s



SCOOP IMPLEMENTATION MODEL 2

Lock contention can be reduced if there is no memory to contend over:
use a queue of queues



Queue of queues: 15s.


With limited dequeue spinning: 3.7s

Don't use the queue for "second" queries: 2.0s



SCOOP IMPLEMENTATION MODEL 2

Lock contention can be reduced if there is no memory to contend over:
use a queue of queues



Queue of queues: 15s.

With limited dequeue spinning: 3.7s

Don't use the queue for "second" queries: 2.0s




PUTTING THEM ALL TOGETHER

Approach	Time
Blocking (STL wrapper)	404s
TBB Non-blocking (sleep)	180s
ES 7.2	155s
TBB Blocking	108s
ES 7.1	33s
Queue of queues	15s
QoQ + Spinning	3.7s
QoQ + Local queries	2.0s




EIFFELSTUDIO 7.1 VS. PROOF OF CONCEPT





NEW BENCHMARK RESULTS





IMPLEMENTATION COMMENTS

- ▶ The processor “runtime” is small: about 150 lines of C++.
- ▶ Seems to be the right target abstraction; almost no usage of explicit synchronization.
- ▶ Only one dependency, TBB.
- ▶ Probably still some room to fine tune, i.e. use specialized queues where possible, like single-producer/consumer queues.